

Progress Report 2
FHIR-Based Prediction of Hospital Readmission for Patients Aged 50 Plus
Group 13
Group Members: Frank Ofosu & Oscar Odera

1. AI Model(s) Used in Research Project

Our project utilizes multiple machine learning models including:

- Random Forest
- XGBoost
- MLP
- Logistic Regression
- Gradient Boosting
- SVM
- KNN
- Adaboost
- Naive Bayesian

Explain the architecture and key components of your model(s)

Out of the models above, Random Forest and Gradient Boosting performed well. Their key architecture and key components include:

A. Random Forest Classifier: Ensemble model based on decision trees, trains multiple trees on different bootstrap samples, uses **majority voting** for classification and tunes parameters with `n_estimators`, `max_depth`

B. Gradient Boosting Classifier: It is a sequential tree-based ensemble model, each tree attempts to correct the errors of its predecessor and it is tuned using `n_estimators`, `learning_rate`

Justify why you selected this particular model/approach for your AI medical/health project.

In predicting 30-day hospital readmission using patient level data, a complex classification task involving multiple structured features like labs, vitals, and chronic condition flags. The Tree based models chosen above are best in handling mixed feature types, capturing non-linear relationships and are best in providing feature importance for explainability which is very critical in healthcare.

2. Performance Metrics Analysis

Sample of Random Forest before Hyperparameter tuning

```
Evaluating: Random Forest
Accuracy: 0.760
F1 Score: 0.623
AUROC: 0.828
Confusion Matrix:
[[54 10]
 [13 19]]
Classification Report:
              precision    recall  f1-score   support

    0.0         0.81      0.84      0.82         64
    1.0         0.66      0.59      0.62         32

   accuracy          0.76
  macro avg          0.73
 weighted avg          0.76
```

Accuracy; the percentage of total correct predictions made by the model, it gives us an idea of how often the model is right.


Precision; the percentage of patients predicted as readmitted who were actually readmitted, it helps us minimize false alarms.

Recall (Sensitivity); the percentage of actual readmitted patients that were correctly identified, it helps us to capture all high risk patients.

F1 Score; Harmonic mean of precision and recall, its useful in balancing the trade-off between precision and recall.

AUC – ROC; Measures the ability of our models to distinguish between readmitted and non-admitted patients, across all thresholds, we use it as a Gold standard in our research.

Performance after hyperparameter tuning for all models



Compare your current results with expected benchmarks or standards in the field (this should stem from the publications you are referencing as relevant literature)

The performance of our models are way better compared to legacy models like LACE and HOSPITAL Scores and also better after hyperparameter tuning than that of Michailidis et al. that obtained an area under the curve (AUC) of 0.78 by utilizing a Random Forest classifier with administrative, clinical-medical, and operational data.

3. Project Status Summary

Our project is on track for completion by April 18th (90% done as seen in the attached project code). Our next steps include preparing presentations and creating Github repositories.

```

import os
import json
import pandas as pd
from tqdm import tqdm

json_folder = r"C:\Users\admin\Desktop\AI in Healthcare\Research
Project\All Converted Files"
observations = []


for filename in tqdm(os.listdir(json_folder)):
    if filename.endswith(".json"):
        path = os.path.join(json_folder, filename)
        try:
            with open(path) as f:
                bundle = json.load(f)
                for entry in bundle.get("entry", []):
                    resource = entry.get("resource", {})
                    if resource.get("resourceType") == "Observation":
                        patient_ref = resource.get("subject",
{}).get("reference", "")
                        patient_id = patient_ref.split("/")[-1] if
patient_ref else None
                        raw_time = resource.get("effectiveDateTime")

                        if patient_id and raw_time:
                            ts = pd.to_datetime(raw_time,
errors='coerce', utc=True)
                            ts_naive = ts.tz_convert(None) if
ts.tzinfo else ts

                            observations.append({
                                "patient_id": patient_id,
                                "timestamp": ts_naive,
                                "file": filename
                            })
                        except Exception as e:
                            print(f"Error reading {filename}: {e}")

```

100%|



```

# Create DataFrame
obs_df = pd.DataFrame(observations)

# Drop rows with missing timestamps
obs_df = obs_df.dropna(subset=["timestamp"])

# Convert patient ID to string (for grouping)
obs_df["patient_id"] = obs_df["patient_id"].astype(str)

```

```
# Sort by patient and timestamp safely
obs_df = obs_df.sort_values(by=["patient_id",
"timestamp"]).reset_index(drop=True)

# Preview
print(f"Total observations extracted: {len(obs_df)}")
obs_df.head()
```

Total observations extracted: 39649

```
patient_id      timestamp \
0  000000001 2024-10-11 01:01:00
1  000000001 2024-10-11 01:01:00
2  000000001 2024-10-11 01:04:01
3  000000001 2024-10-11 03:35:00
4  000000001 2024-10-11 03:35:00
```

```
file
0  fhir_output_TEST_ccd_uphealthsystem-marquette_...
1  fhir_output_TEST_ccd_uphealthsystem-marquette_...
2  fhir_output_TEST_ccd_uphealthsystem-marquette_...
3  fhir_output_TEST_ccd_uphealthsystem-marquette_...
4  fhir_output_TEST_ccd_uphealthsystem-marquette_...
```

```
# Start with no readmissions
obs_df["readmitted"] = 0

# Loop through observations by same patient
for i in range(1, len(obs_df)):
    prev = obs_df.iloc[i - 1]
    curr = obs_df.iloc[i]

    if curr["patient_id"] == prev["patient_id"]:
        days_apart = (curr["timestamp"] - prev["timestamp"]).days
        if 0 < days_apart <= 30:
            obs_df.at[i, "readmitted"] = 1

# If patient was readmitted at least once → label = 1
patient_labels = obs_df.groupby("patient_id")
["readmitted"].max().reset_index()

# Convert patient_id to string to match feature dataset
patient_labels["patient_id"] =
patient_labels["patient_id"].astype(str)

# Preview the readmission label distribution
print("Readmission Label Counts:")
print(patient_labels["readmitted"].value_counts())

patient_labels.head()
```

Readmission Label Counts:

readmitted

0 200

1 152

Name: count, dtype: int64

	patient_id	readmitted
0	000000001	1
1	000000003	1
2	000000004	0
3	000000005	1
4	000000007	1

```
import pandas as pd
```

```
# Load your modeling dataset
```

```
df_main = pd.read_csv("readmission_dataset2.csv")
```

```
# Ensure both patient IDs are strings for a safe merge
```

```
df_main["patient_id"] = df_main["patient_id"].astype(str)
```

```
patient_labels["patient_id"] =
```

```
patient_labels["patient_id"].astype(str)
```

```
# Merge real readmission labels
```

```
df_labeled = df_main.merge(patient_labels, on="patient_id",  
how="left")
```

```
# Patients not found in readmission label → assume not readmitted
```

```
df_labeled["readmitted"] =
```

```
df_labeled["readmitted"].fillna(0).astype(int)
```

```
# Confirm success
```

```
print("Readmission label counts after merge:")
```

```
print(df_labeled["readmitted"].value_counts())
```

```
df_labeled.head()
```

Readmission label counts after merge:

readmitted

0 320

1 158

Name: count, dtype: int64

	patient_id	gender	age	num_conditions	num_encounters	\
0	106	female	61.0	4	2	
1	109	female	80.0	36	4	
2	107	male	39.0	10	1	
3	212	NaN	NaN	2	1	
4	148	male	27.0	15	2	

	chronic_conditions	Height	Glucose_avg	Temperature
SystolicBP	\			

0		NaN	129.666667	156.0	36.80
115.000000					
1		NaN	157.500000	NaN	36.86
131.250000					
2		NaN	162.600000	NaN	36.56
131.000000					
3	Asthma; Hypertension		190.000000	NaN	37.20
175.000000					
4	Depression		173.600000	NaN	36.97
112.571429					

	DiastolicBP	HeartRate	HemoglobinA1C	Weight	BMI
readmitted					
0	76.666667	99.333333	11.6	NaN	NaN
0					
1	92.250000	94.000000	NaN	48.762	19.6625
1					
2	84.000000	114.000000	NaN	65.318	24.7200
0					
3	91.000000	NaN	NaN	NaN	NaN
0					
4	63.857143	65.142857	NaN	115.667	39.9400
0					

Drop text columns not used in modeling (optional: add more if needed)

```
df_model = df_labeled.drop(columns=["patient_id", "birthDate",
"num_encounters", "address", "name"], errors='ignore')
```

Map gender manually

```
if "gender" in df_model.columns:
    df_model["gender"] = df_model["gender"].map({"female": 1, "male":
0})
```

Automatically one-hot encode remaining object columns

```
non_numeric =
df_model.select_dtypes(include='object').columns.tolist()
if non_numeric:
    print("One-hot encoding:", non_numeric)
    df_model = pd.get_dummies(df_model, columns=non_numeric,
drop_first=True)
```

One-hot encoding: ['chronic_conditions']

```
from sklearn.impute import SimpleImputer
```

Use SimpleImputer to fill remaining NaNs

```
imputer = SimpleImputer(strategy="median")
df_model_imputed = pd.DataFrame(
    imputer.fit_transform(df_model),
```

```

        columns=df_model.columns
    )

    # Confirm no NaNs remain
    print("NaNs after imputation:", df_model_imputed.isnull().sum().sum())

    NaNs after imputation: 0

    from sklearn.model_selection import train_test_split

    X = df_model_imputed.drop(columns=["readmitted"])
    y = df_model_imputed["readmitted"]

    X_train, X_test, y_train, y_test = train_test_split(
        X, y, stratify=y, test_size=0.2, random_state=42
    )

    from sklearn.utils import resample
    import numpy as np

    # Combine for resampling
    train_df = pd.concat([X_train, y_train], axis=1)

    # Split by class
    majority = train_df[train_df["readmitted"] == 0]
    minority = train_df[train_df["readmitted"] == 1]

    # Upsample minority to match majority
    minority_upsampled = resample(minority, replace=True,
                                  n_samples=len(majority), random_state=42)

    # Combine and shuffle
    balanced_df = pd.concat([majority, minority_upsampled]).sample(frac=1,
                                                                    random_state=42)

    # Separate again
    X_res = balanced_df.drop(columns=["readmitted"])
    y_res = balanced_df["readmitted"]

    print("Balanced class distribution:", np.bincount(y_res))

    Balanced class distribution: [256 256]

    from sklearn.ensemble import (
        RandomForestClassifier, GradientBoostingClassifier,
        AdaBoostClassifier
    )
    from sklearn.linear_model import LogisticRegression
    from sklearn.svm import SVC
    from sklearn.neighbors import KNeighborsClassifier
    from sklearn.naive_bayes import GaussianNB

```



```

from sklearn.neural_network import MLPClassifier

# Optional: try XGBoost if installed
try:
    from xgboost import XGBClassifier
except:
    XGBClassifier = None

# Define models dictionary
models = {
    "Random Forest": RandomForestClassifier(random_state=42),
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Gradient Boosting": GradientBoostingClassifier(),
    "AdaBoost": AdaBoostClassifier(),
    "SVM": SVC(probability=True),
    "KNN": KNeighborsClassifier(),
    "Naive Bayes": GaussianNB(),
    "MLP Neural Net": MLPClassifier(max_iter=1000)
}

if XGBClassifier:
    models["XGBoost"] = XGBClassifier(eval_metric='logloss',
enable_categorical=False)

import warnings
warnings.filterwarnings("ignore")

from sklearn.metrics import (
    accuracy_score, f1_score, roc_auc_score, confusion_matrix,
    classification_report
)

results = []

for name, model in models.items():
    print(f"\nEvaluating: {name}")
    try:
        model.fit(X_res, y_res)
        y_pred = model.predict(X_test)

        # Use predicted probabilities for AUC if available
        if hasattr(model, "predict_proba"):
            y_proba = model.predict_proba(X_test)[:, 1]
        else:
            y_proba = y_pred

        acc = accuracy_score(y_test, y_pred)
        f1 = f1_score(y_test, y_pred, zero_division=0)
        auc = roc_auc_score(y_test, y_proba)

```

```

        print(f"Accuracy: {acc:.3f}")
        print(f"F1 Score: {f1:.3f}")
        print(f"AUROC: {auc:.3f}")
        print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
        print("Classification Report:\n",
classification_report(y_test, y_pred, zero_division=0))

        results.append({"Model": name, "Accuracy": acc, "F1 Score":
f1, "AUROC": auc})

    except Exception as e:
        print(f"{name} failed: {e}")

```

Evaluating: Random Forest

Accuracy: 0.760

F1 Score: 0.623

AUROC: 0.828

Confusion Matrix:

[[54 10]

[13 19]]

Classification Report:

	precision	recall	f1-score	support
0.0	0.81	0.84	0.82	64
1.0	0.66	0.59	0.62	32
accuracy			0.76	96
macro avg	0.73	0.72	0.72	96
weighted avg	0.76	0.76	0.76	96

Evaluating: Logistic Regression

Accuracy: 0.656

F1 Score: 0.459

AUROC: 0.648

Confusion Matrix:

[[49 15]

[18 14]]

Classification Report:

	precision	recall	f1-score	support
0.0	0.73	0.77	0.75	64
1.0	0.48	0.44	0.46	32
accuracy			0.66	96
macro avg	0.61	0.60	0.60	96
weighted avg	0.65	0.66	0.65	96

Evaluating: Gradient Boosting

Accuracy: 0.729

F1 Score: 0.594

AUROC: 0.847

Confusion Matrix:

[[51 13]

[13 19]]

Classification Report:

	precision	recall	f1-score	support
0.0	0.80	0.80	0.80	64
1.0	0.59	0.59	0.59	32
accuracy			0.73	96
macro avg	0.70	0.70	0.70	96
weighted avg	0.73	0.73	0.73	96

Evaluating: AdaBoost

Accuracy: 0.760

F1 Score: 0.667

AUROC: 0.844

Confusion Matrix:

[[50 14]

[9 23]]

Classification Report:

	precision	recall	f1-score	support
0.0	0.85	0.78	0.81	64
1.0	0.62	0.72	0.67	32
accuracy			0.76	96
macro avg	0.73	0.75	0.74	96
weighted avg	0.77	0.76	0.76	96

Evaluating: SVM

Accuracy: 0.740

F1 Score: 0.468

AUROC: 0.720

Confusion Matrix:

[[60 4]

[21 11]]

Classification Report:

	precision	recall	f1-score	support
0.0	0.74	0.94	0.83	64
1.0	0.73	0.34	0.47	32
accuracy			0.74	96

macro avg	0.74	0.64	0.65	96
weighted avg	0.74	0.74	0.71	96

Evaluating: KNN

Accuracy: 0.677

F1 Score: 0.551

AUROC: 0.740

Confusion Matrix:

[[46 18]

[13 19]]

Classification Report:

	precision	recall	f1-score	support
0.0	0.78	0.72	0.75	64
1.0	0.51	0.59	0.55	32

accuracy			0.68	96
macro avg	0.65	0.66	0.65	96
weighted avg	0.69	0.68	0.68	96

Evaluating: Naive Bayes

Accuracy: 0.552

F1 Score: 0.583

AUROC: 0.787

Confusion Matrix:

[[23 41]

[2 30]]

Classification Report:

	precision	recall	f1-score	support
0.0	0.92	0.36	0.52	64
1.0	0.42	0.94	0.58	32

accuracy			0.55	96
macro avg	0.67	0.65	0.55	96
weighted avg	0.75	0.55	0.54	96

Evaluating: MLP Neural Net

Accuracy: 0.729

F1 Score: 0.536

AUROC: 0.774

Confusion Matrix:

[[55 9]

[17 15]]

Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

	0.0	0.76	0.86	0.81	64
	1.0	0.62	0.47	0.54	32
accuracy				0.73	96
macro avg		0.69	0.66	0.67	96
weighted avg		0.72	0.73	0.72	96

Evaluating: XGBoost

Accuracy: 0.719

F1 Score: 0.571

AUROC: 0.808

Confusion Matrix:

```
[[51 13]
```

```
[14 18]]
```

Classification Report:

	precision	recall	f1-score	support
0.0	0.78	0.80	0.79	64
1.0	0.58	0.56	0.57	32
accuracy			0.72	96
macro avg	0.68	0.68	0.68	96
weighted avg	0.72	0.72	0.72	96

```
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import f1_score, recall_score
# Define models + hyperparameter grids
model_configs = {
    "Logistic Regression": (
        LogisticRegression(max_iter=1000),
        {"C": [0.1, 1, 10], "solver": ["liblinear", "lbfgs"]}
    ),
    "Random Forest": (
        RandomForestClassifier(),
        {"n_estimators": [100, 200], "max_depth": [5, 10]}
    ),
    "XGBoost": (
        XGBClassifier(use_label_encoder=False, eval_metric='logloss'),
        {"n_estimators": [100, 200], "max_depth": [3, 6]}
    ),
    "SVM": (
        SVC(probability=True),
        {"C": [0.1, 1], "kernel": ["linear", "rbf"]}
    ),
    "KNN": (
        KNeighborsClassifier(),
        {"n_neighbors": [3, 5, 7]}
    ),
}
```

```

    "Naive Bayes": (
        GaussianNB(),
        {} # No tuning
    ),
    "Gradient Boosting": (
        GradientBoostingClassifier(),
        {"n_estimators": [100, 200], "max_depth": [3, 6]}
    ),
    "MLP": (
        MLPClassifier(max_iter=500),
        {"hidden_layer_sizes": [(50,), (100,)], "activation": ["relu",
"tanh"]}
    )
}

# Tune + evaluate
models = {}
results = {}

for name, (model, params) in model_configs.items():
    print(f"Tuning {name}...")
    clf = GridSearchCV(model, params, scoring="roc_auc", cv=3,
n_jobs=-1)
    clf.fit(X_train, y_train)
    best = clf.best_estimator_
    models[name] = best

    y_pred = best.predict(X_test)
    y_prob = best.predict_proba(X_test)[:, 1]

    results[name] = {
        "F1": f1_score(y_test, y_pred),
        "Recall": recall_score(y_test, y_pred),
        "AUROC": roc_auc_score(y_test, y_prob)
    }

Tuning Logistic Regression...
Tuning Random Forest...
Tuning XGBoost...
Tuning SVM...
Tuning KNN...
Tuning Naive Bayes...
Tuning Gradient Boosting...
Tuning MLP...

import pandas as pd

# Create summary table (transpose first)
results_df = pd.DataFrame(results).T.sort_values(by="AUROC",
ascending=False).reset_index()

```

```
# Rename for clarity
results_df.rename(columns={'index': 'Model'}, inplace=True)
```


```
# Display
print("Final Model Comparison:")
display(results_df)
```

Final Model Comparison:

	Model	F1	Recall	AUROC
0	Gradient Boosting	0.688525	0.65625	0.870117
1	Random Forest	0.677419	0.65625	0.861816
2	XGBoost	0.593750	0.59375	0.805176
3	MLP	0.590164	0.56250	0.798828
4	SVM	0.222222	0.12500	0.796875
5	Naive Bayes	0.520000	0.40625	0.776855
6	KNN	0.500000	0.43750	0.763184
7	Logistic Regression	0.390244	0.25000	0.708496

```
import seaborn as sns
import matplotlib.pyplot as plt
print("\nConfusion Matrix for Random Forest:")
cm = confusion_matrix(y_test, models["Random Forest"].predict(X_test))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=[0,1],
yticklabels=[0,1])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Random Forest Confusion Matrix")
plt.show()
```

Confusion Matrix for Random Forest:



```
import matplotlib.pyplot as plt
from sklearn.metrics import roc_curve, auc

plt.figure(figsize=(10, 6))


for name, model in models.items():
    try:
        if hasattr(model, "predict_proba"):
            y_score = model.predict_proba(X_test)[:, 1]
        else:
            y_score = model.decision_function(X_test)

        fpr, tpr, _ = roc_curve(y_test, y_score)
        roc_auc = auc(fpr, tpr)
        plt.plot(fpr, tpr, label=f"{name} (AUC = {roc_auc:.2f})")
    except:
        continue

plt.plot([0, 1], [0, 1], "k--", label="Random Guess")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curves for All Models")
plt.legend(loc="lower right")
```



```
plt.grid()  
plt.show()
```



```
# Must match the model that supports .feature_importances_  
rf_model = models["Random Forest"]  
importances = rf_model.feature_importances_  
  
# Create a Series for plotting  
feat_imp = pd.Series(importances,  
index=X.columns).sort_values(ascending=False)  
  
# Plot top 10 features  
feat_imp.head(10).plot(kind="barh", figsize=(8, 6), title="Top 10  
Feature Importances")  
plt.gca().invert_yaxis()  
plt.xlabel("Importance Score")  
plt.tight_layout()  
plt.show()
```

